

Basic Data Exploration and Prediction

```
setwd("C:/Users/baoji/OneDrive/IVLE/ST4240 Data Mining_2017/Predictive Modelling Assignment")

set.seed(1) #for reproducibility
library(plyr)
library(dplyr, warn.conflicts = FALSE)
library(ggplot2)
library(corrplot)
library(e1071)
library(hexbin)
library(gbm)


## Loading required package: survival
## Loading required package: lattice
## Loading required package: splines
## Loading required package: parallel
## Loaded gbm 2.1.1
library(xgboost)

##
## Attaching package: 'xgboost'
## The following object is masked from 'package:dplyr':
##
##      slice
library(lars)


## Loaded lars 1.2
library(Matrix)

filename_train = "train.csv"
data = read.csv(file = filename_train)


data %>%
  ggplot() + geom_point(aes(x=x, y=y, colour=price/living_area),alpha=0.3) +
  scale_colour_gradient(low = "blue", high="red")
```




```
data %>%  
  ggplot(aes(x=factor.bedrooms, y=price)) +  
  geom_boxplot()
```



```
data %>%  
  ggplot(aes(x=factor(bathrooms), y=living_area)) +  
  geom_boxplot()
```

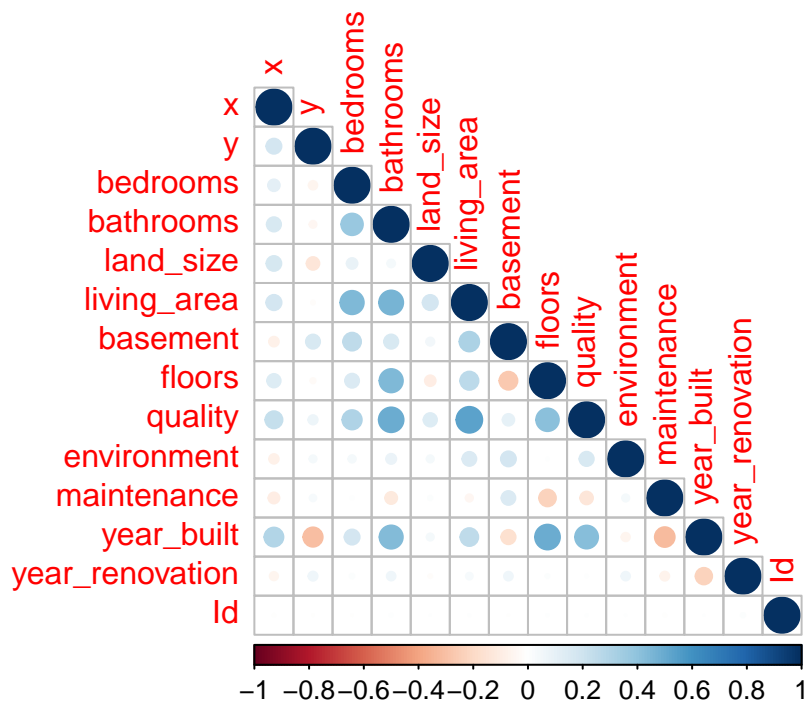


```
data %>%  
  ggplot(aes(x=factor(floors), y=price)) +  
  geom_boxplot()
```




```
# data_skewneww_score <- sapply(data, skewness)
# data_skewneww_score

correlations<- cor(data[,-1],use="everything")
corrplot(correlations, method="circle", type="lower", sig.level = 0.01, insig = "blank")
```



```
# ggplot(data,aes(x= year_built,y=price))+geom_point()+geom_smooth()
# ggplot(data,aes(x= year_renovation,y=price))+geom_point()+geom_smooth()

data %>%
  ggplot(aes(x=living_area,y=price)) +
  stat_binhex(bins=50)
```



Data Preprocessing

```
x_data <- data
x_data$price <- NULL
y_data <- data$price

n = nrow(data)

filename_output = "test.csv"
data_output = read.csv(file = filename_output)
output_id = data_output$Id

ALL_X_DATA = rbind(x_data,data_output)

renovation_year_generator <- function(x) {
  if (x['year_renovation'] == -1){
    return (x['year_built'])
  }
  else{
    return (2017-x['year_renovation'])
  }
}

ETL <- function(data){

  # change year built to age of the house
  data$year_built = 2017-data$year_built
```

```

# a dummy variable which indicates if the house has been renovated
data$is_renovated = as.numeric(data$year_renovation != -1)

# years since last renovation
data$year_renovation = apply(data, 1, renovation_year_generator)

# indicate if the house as a basement
data$has_basement = as.numeric(data$basement != 0)

# area per floor
data$area_per_floor = data$living_area/data$floors

# indicate the house is old (older than 75 years)
data$is_old = as.numeric(data$year_built >= 75)

# taking square of the scores
data$quality_2 = data$quality^2
data$environment_2 = data$environment^2
data$maintenance_2 = data$maintenance^2
data <- data[, !(names(data) %in% c("Id"))]

#clustering based on location (15 areas)
location_data <- data[,c('x','y')]
set.seed(1) #for reproducibility
location_fit <- kmeans(location_data, 15)
location.f = factor(location_fit$cluster)
location_dummies = model.matrix(~location.f)[-1]
data <- cbind(data,location_dummies)

#clustering based on room type
type_data <- data[,c('bedrooms','bathrooms','living_area','floors','has_basement','area_per_floor')]
set.seed(1) #for reproducibility
type_fit <- kmeans(type_data, 3)
type.f = factor(type_fit$cluster)
type_dummies = model.matrix(~type.f)[-1]
data <- cbind(data,type_dummies)

return (data)
}

ALL_X_DATA <- ETL(ALL_X_DATA)
x_data = ALL_X_DATA[1:n,]
data <- x_data
data$price <- y_data

data_output = ALL_X_DATA[(n+1):nrow(ALL_X_DATA),]

# create a sample vector of test values
set.seed(1) #for reproducibility
test.n <- sample(1:nrow(data), nrow(data)/5, replace = F)

# test dataset
y_test <- y_data[test.n]

```



```

x_test <- x_data[test.n,]
test <- x_test
test$price <- y_test
# x_test <- test[, !(names(test) %in% c("price"))]

# test <- data
# y_test <- test$price
# x_test <- test[, !(names(cleaned_data) %in% c("price"))]

# train dataset
y_train <- y_data[-test.n]
x_train <- x_data[-test.n,]
train <- x_train
train$price <- y_train

RMSPE <- function(pred_y,true_y){
  a <- sqrt(mean(((pred_y-true_y)/true_y)^2))
  return(a)
}

result_output <- function(y){
  # y <- expm1(y)
  y <- round_any(y,100)
  return(y)
}

xg_eval_mae <- function (yhat, dtrain) {
  y = getinfo(dtrain, "label")
  y = expm1(y)
  yhat = expm1(yhat)
  err= RMSPE(yhat, y)
  return (list(metric = "error", value = err))
}

```

linear regression (stepwise) Since prices of the house are skewed, we take the log the price to reduce skewness.

```

linear_model = lm(log1p(price) ~ ., data = train)
linear_model <- step(linear_model)

```

```

## Start:  AIC=-17546.5
## log1p(price) ~ x + y + bedrooms + bathrooms + land_size + living_area +
##   basement + floors + quality + environment + maintenance +
##   year_built + year_renovation + is_renovated + has_basement +
##   area_per_floor + is_old + quality_2 + environment_2 + maintenance_2 +
##   location.f2 + location.f3 + location.f4 + location.f5 + location.f6 +
##   location.f7 + location.f8 + location.f9 + location.f10 +
##   location.f11 + location.f12 + location.f13 + location.f14 +
##   location.f15 + type.f2 + type.f3
##
##           Df Sum of Sq  RSS   AIC
## - location.f11      1    0.007 884.16 -17548
## - floors            1    0.012 884.17 -17548
## - location.f6       1    0.044 884.20 -17548
## - environment       1    0.074 884.23 -17548
## - maintenance       1    0.074 884.23 -17548

```

```

## <none> 884.16 -17547
## - area_per_floor 1 0.249 884.41 -17546
## - basement 1 0.388 884.54 -17545
## - location.f14 1 0.424 884.58 -17545
## - year_renovation 1 0.592 884.75 -17543
## - year_built 1 0.785 884.94 -17541
## - has_basement 1 1.029 885.19 -17539
## - type.f3 1 1.207 885.36 -17538
## - maintenance_2 1 1.552 885.71 -17535
## - is_renovated 1 1.667 885.82 -17533
## - location.f9 1 2.625 886.78 -17525
## - x 1 2.645 886.80 -17525
## - location.f13 1 2.996 887.15 -17521
## - type.f2 1 3.123 887.28 -17520
## - environment_2 1 3.328 887.48 -17518
## - location.f12 1 3.464 887.62 -17517
## - quality_2 1 4.489 888.65 -17508
## - location.f8 1 5.471 889.63 -17499
## - is_old 1 5.569 889.73 -17498
## - location.f7 1 6.168 890.32 -17493
## - bathrooms 1 8.375 892.53 -17473
## - location.f10 1 9.627 893.78 -17462
## - land_size 1 10.114 894.27 -17458
## - location.f4 1 10.298 894.45 -17456
## - living_area 1 13.170 897.33 -17430
## - location.f2 1 15.846 900.00 -17406
## - bedrooms 1 18.954 903.11 -17379
## - y 1 21.425 905.58 -17357
## - location.f3 1 27.189 911.35 -17306
## - location.f15 1 28.420 912.58 -17295
## - location.f5 1 33.840 918.00 -17248
## - quality 1 34.882 919.04 -17239
##
## Step: AIC=-17548.44
## log1p(price) ~ x + y + bedrooms + bathrooms + land_size + living_area +
## basement + floors + quality + environment + maintenance +
## year_built + year_renovation + is_renovated + has_basement +
## area_per_floor + is_old + quality_2 + environment_2 + maintenance_2 +
## location.f2 + location.f3 + location.f4 + location.f5 + location.f6 +
## location.f7 + location.f8 + location.f9 + location.f10 +
## location.f12 + location.f13 + location.f14 + location.f15 +
## type.f2 + type.f3
##
## Df Sum of Sq RSS AIC
## - floors 1 0.012 884.18 -17550
## - environment 1 0.074 884.24 -17550
## - maintenance 1 0.075 884.24 -17550
## - location.f6 1 0.076 884.24 -17550
## <none> 884.16 -17548
## - area_per_floor 1 0.250 884.41 -17548
## - basement 1 0.387 884.55 -17547
## - year_renovation 1 0.592 884.76 -17545
## - location.f14 1 0.687 884.85 -17544
## - year_built 1 0.785 884.95 -17543

```

```

## - has_basement      1      1.041 885.20 -17541
## - type.f3           1      1.213 885.38 -17540
## - maintenance_2    1      1.552 885.72 -17536
## - is_renovated      1      1.667 885.83 -17535
## - type.f2           1      3.134 887.30 -17522
## - environment_2     1      3.330 887.49 -17520
## - x                 1      3.651 887.81 -17518
## - quality_2         1      4.489 888.65 -17510
## - location.f9        1      5.035 889.20 -17505
## - location.f12       1      5.454 889.62 -17501
## - is_old            1      5.678 889.84 -17499
## - location.f13       1      5.743 889.91 -17499
## - location.f8        1      7.800 891.96 -17480
## - bathrooms         1      8.378 892.54 -17475
## - land_size         1     10.108 894.27 -17460
## - location.f4        1     10.593 894.76 -17455
## - location.f7        1     10.666 894.83 -17455
## - location.f10       1     12.391 896.56 -17439
## - living_area        1     13.182 897.35 -17432
## - bedrooms          1     18.948 903.11 -17381
## - y                 1     24.987 909.15 -17328
## - location.f2        1     27.483 911.65 -17306
## - location.f3        1     28.861 913.03 -17294
## - quality            1     34.876 919.04 -17241
## - location.f15       1     35.868 920.03 -17232
## - location.f5        1     53.077 937.24 -17084
##
## Step: AIC=-17550.33
## log1p(price) ~ x + y + bedrooms + bathrooms + land_size + living_area +
##   basement + quality + environment + maintenance + year_built +
##   year_renovation + is_renovated + has_basement + area_per_floor +
##   is_old + quality_2 + environment_2 + maintenance_2 + location.f2 +
##   location.f3 + location.f4 + location.f5 + location.f6 + location.f7 +
##   location.f8 + location.f9 + location.f10 + location.f12 +
##   location.f13 + location.f14 + location.f15 + type.f2 + type.f3
##
##           Df Sum of Sq   RSS   AIC
## - location.f6      1      0.073 884.25 -17552
## - maintenance      1      0.074 884.25 -17552
## - environment       1      0.075 884.25 -17552
## <none>                                884.18 -17550
## - basement         1      0.378 884.55 -17549
## - area_per_floor    1      0.412 884.59 -17549
## - year_renovation   1      0.592 884.77 -17547
## - location.f14      1      0.697 884.87 -17546
## - year_built        1      0.774 884.95 -17545
## - has_basement      1      1.081 885.26 -17543
## - type.f3           1      1.206 885.38 -17541
## - maintenance_2    1      1.546 885.72 -17538
## - is_renovated      1      1.661 885.84 -17537
## - type.f2           1      3.124 887.30 -17524
## - environment_2     1      3.319 887.49 -17522
## - x                 1      3.717 887.89 -17519
## - quality_2         1      4.533 888.71 -17511

```

```

## - location.f9      1      5.065 889.24 -17507
## - location.f12     1      5.472 889.65 -17503
## - is_old           1      5.687 889.86 -17501
## - location.f13     1      5.739 889.91 -17501
## - location.f8      1      7.836 892.01 -17482
## - bathrooms       1      8.372 892.55 -17477
## - land_size        1     10.103 894.28 -17461
## - location.f4      1     10.590 894.77 -17457
## - location.f7      1     10.657 894.83 -17457
## - location.f10     1     12.391 896.57 -17441
## - bedrooms        1     18.939 903.11 -17383
## - living_area      1     22.551 906.73 -17351
## - y               1     25.045 909.22 -17329
## - location.f2      1     27.474 911.65 -17308
## - location.f3      1     28.898 913.07 -17295
## - quality          1     35.551 919.73 -17237
## - location.f15     1     35.857 920.03 -17234
## - location.f5      1     53.113 937.29 -17086
##
## Step: AIC=-17551.67
## log1p(price) ~ x + y + bedrooms + bathrooms + land_size + living_area +
##   basement + quality + environment + maintenance + year_built +
##   year_renovation + is_renovated + has_basement + area_per_floor +
##   is_old + quality_2 + environment_2 + maintenance_2 + location.f2 +
##   location.f3 + location.f4 + location.f5 + location.f7 + location.f8 +
##   location.f9 + location.f10 + location.f12 + location.f13 +
##   location.f14 + location.f15 + type.f2 + type.f3
##
##           Df Sum of Sq  RSS   AIC
## - maintenance      1      0.072 884.32 -17553
## - environment       1      0.075 884.32 -17553
## <none>                                884.25 -17552
## - basement          1      0.370 884.62 -17550
## - area_per_floor    1      0.415 884.66 -17550
## - year_renovation   1      0.591 884.84 -17548
## - location.f14      1      0.688 884.94 -17548
## - year_built        1      0.772 885.02 -17547
## - has_basement      1      1.085 885.33 -17544
## - type.f3           1      1.220 885.47 -17543
## - maintenance_2     1      1.537 885.79 -17540
## - is_renovated      1      1.659 885.91 -17539
## - type.f2           1      3.138 887.39 -17525
## - environment_2     1      3.315 887.56 -17524
## - x                 1      3.893 888.14 -17519
## - quality_2         1      4.589 888.84 -17512
## - location.f9       1      5.326 889.57 -17506
## - location.f12      1      5.406 889.65 -17505
## - is_old            1      5.877 890.13 -17501
## - location.f13      1      6.404 890.65 -17496
## - bathrooms         1      8.396 892.64 -17478
## - land_size         1     10.081 894.33 -17463
## - location.f4       1     10.829 895.08 -17456
## - location.f8       1     11.131 895.38 -17454
## - location.f10      1     12.448 896.70 -17442

```

```

## - location.f7      1      12.596 896.84 -17441
## - bedrooms        1      18.911 903.16 -17384
## - living_area     1      22.596 906.84 -17352
## - location.f2     1      27.415 911.66 -17309
## - location.f3     1      28.987 913.23 -17296
## - y               1      29.354 913.60 -17292
## - quality         1      35.679 919.93 -17237
## - location.f15    1      36.300 920.55 -17232
## - location.f5     1      68.299 952.55 -16959
##
## Step:  AIC=-17553.02
## log1p(price) ~ x + y + bedrooms + bathrooms + land_size + living_area +
##   basement + quality + environment + year_built + year_renovation +
##   is_renovated + has_basement + area_per_floor + is_old + quality_2 +
##   environment_2 + maintenance_2 + location.f2 + location.f3 +
##   location.f4 + location.f5 + location.f7 + location.f8 + location.f9 +
##   location.f10 + location.f12 + location.f13 + location.f14 +
##   location.f15 + type.f2 + type.f3
##
##           Df Sum of Sq   RSS   AIC
## - environment      1      0.075 884.40 -17554
## <none>                884.32 -17553
## - basement         1      0.371 884.69 -17552
## - area_per_floor   1      0.419 884.74 -17551
## - year_renovation  1      0.589 884.91 -17550
## - location.f14     1      0.680 885.00 -17549
## - year_built       1      0.777 885.10 -17548
## - has_basement     1      1.080 885.40 -17545
## - type.f3          1      1.217 885.54 -17544
## - is_renovated     1      1.664 885.98 -17540
## - type.f2          1      3.139 887.46 -17527
## - environment_2    1      3.319 887.64 -17525
## - x                1      3.895 888.21 -17520
## - quality_2        1      4.579 888.90 -17514
## - location.f9      1      5.304 889.62 -17507
## - location.f12     1      5.375 889.69 -17507
## - is_old           1      5.945 890.27 -17501
## - location.f13     1      6.386 890.71 -17498
## - bathrooms        1      8.413 892.73 -17479
## - maintenance_2    1      8.429 892.75 -17479
## - land_size        1     10.065 894.38 -17465
## - location.f4      1     10.793 895.11 -17458
## - location.f8      1     11.147 895.47 -17455
## - location.f10     1     12.411 896.73 -17444
## - location.f7      1     12.604 896.92 -17442
## - bedrooms         1     18.935 903.26 -17386
## - living_area      1     22.603 906.92 -17353
## - location.f2      1     27.366 911.69 -17311
## - location.f3      1     28.950 913.27 -17297
## - y               1     29.488 913.81 -17293
## - quality          1     35.645 919.97 -17239
## - location.f15     1     36.235 920.56 -17234
## - location.f5      1     68.312 952.63 -16960
##

```

```
## Step: AIC=-17554.34
## log1p(price) ~ x + y + bedrooms + bathrooms + land_size + living_area +
##   basement + quality + year_built + year_renovation + is_renovated +
##   has_basement + area_per_floor + is_old + quality_2 + environment_2 +
##   maintenance_2 + location.f2 + location.f3 + location.f4 +
##   location.f5 + location.f7 + location.f8 + location.f9 + location.f10 +
##   location.f12 + location.f13 + location.f14 + location.f15 +
##   type.f2 + type.f3
##
##           Df Sum of Sq   RSS   AIC
## <none>                884.40 -17554
## - basement           1     0.367 884.76 -17553
## - area_per_floor      1     0.429 884.82 -17553
## - year_renovation     1     0.583 884.98 -17551
## - location.f14        1     0.678 885.07 -17550
## - year_built          1     0.770 885.17 -17549
## - has_basement        1     1.086 885.48 -17547
## - type.f3             1     1.214 885.61 -17545
## - is_renovated        1     1.654 886.05 -17541
## - type.f2             1     3.145 887.54 -17528
## - x                   1     3.890 888.29 -17521
## - quality_2           1     4.612 889.01 -17515
## - location.f9         1     5.293 889.69 -17509
## - location.f12        1     5.367 889.76 -17508
## - is_old              1     5.941 890.34 -17503
## - location.f13        1     6.373 890.77 -17499
## - bathrooms           1     8.405 892.80 -17481
## - maintenance_2       1     8.439 892.83 -17480
## - land_size           1    10.091 894.49 -17466
## - location.f4         1    10.795 895.19 -17459
## - location.f8         1    11.130 895.53 -17456
## - location.f10        1    12.402 896.80 -17445
## - location.f7         1    12.603 897.00 -17443
## - bedrooms            1    18.915 903.31 -17387
## - living_area         1    22.652 907.05 -17354
## - environment_2       1    24.026 908.42 -17342
## - location.f2         1    27.343 911.74 -17313
## - location.f3         1    28.951 913.35 -17299
## - y                   1    29.506 913.90 -17294
## - quality             1    35.728 920.12 -17240
## - location.f15        1    36.184 920.58 -17236
## - location.f5         1    68.327 952.72 -16961
```

```
summary(linear_model)
```

```
##
## Call:
## lm(formula = log1p(price) ~ x + y + bedrooms + bathrooms + land_size +
##   living_area + basement + quality + year_built + year_renovation +
##   is_renovated + has_basement + area_per_floor + is_old + quality_2 +
##   environment_2 + maintenance_2 + location.f2 + location.f3 +
##   location.f4 + location.f5 + location.f7 + location.f8 + location.f9 +
##   location.f10 + location.f12 + location.f13 + location.f14 +
##   location.f15 + type.f2 + type.f3, data = train)
##
```

```

## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.89594 -0.18539  0.01287  0.20356  1.55089
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   1.239e+01  5.320e-02 232.950 < 2e-16 ***
## x             1.502e-01  2.538e-02   5.920 3.36e-09 ***
## y             3.983e-01  2.443e-02  16.305 < 2e-16 ***
## bedrooms      6.536e-02  5.007e-03  13.054 < 2e-16 ***
## bathrooms     6.178e-02  7.099e-03   8.702 < 2e-16 ***
## land_size     3.143e-07  3.296e-08   9.535 < 2e-16 ***
## living_area   1.294e-05  9.059e-07  14.286 < 2e-16 ***
## basement      2.577e-06  1.416e-06   1.819 0.068931 .
## quality       1.948e+00  1.086e-01  17.941 < 2e-16 ***
## year_built    -3.183e-03  1.208e-03  -2.634 0.008445 **
## year_renovation 2.689e-03  1.173e-03   2.293 0.021884 *
## is_renovated   2.742e-01  7.102e-02   3.860 0.000114 ***
## has_basement   4.046e-02  1.294e-02   3.128 0.001769 **
## area_per_floor -1.933e-06  9.835e-07  -1.965 0.049395 *
## is_old         1.215e-01  1.661e-02   7.316 2.80e-13 ***
## quality_2      -7.030e-01  1.091e-01  -6.446 1.22e-10 ***
## environment_2  5.655e-01  3.844e-02  14.713 < 2e-16 ***
## maintenance_2  1.896e-01  2.174e-02   8.720 < 2e-16 ***
## location.f2    -4.195e-01  2.672e-02 -15.695 < 2e-16 ***
## location.f3    -5.671e-01  3.512e-02 -16.150 < 2e-16 ***
## location.f4    -2.307e-01  2.339e-02  -9.862 < 2e-16 ***
## location.f5    -5.217e-01  2.103e-02 -24.811 < 2e-16 ***
## location.f7    -2.578e-01  2.419e-02 -10.656 < 2e-16 ***
## location.f8    -1.697e-01  1.694e-02 -10.014 < 2e-16 ***
## location.f9    -2.431e-01  3.521e-02  -6.906 5.38e-12 ***
## location.f10   -4.279e-01  4.048e-02 -10.571 < 2e-16 ***
## location.f12   -2.928e-01  4.210e-02  -6.954 3.84e-12 ***
## location.f13   -2.143e-01  2.828e-02  -7.577 3.92e-14 ***
## location.f14   -1.138e-01  4.605e-02  -2.471 0.013478 *
## location.f15   -4.068e-01  2.253e-02 -18.056 < 2e-16 ***
## type.f2        9.489e-02  1.782e-02   5.323 1.05e-07 ***
## type.f3        8.578e-02  2.594e-02   3.307 0.000946 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.3332 on 7968 degrees of freedom
## Multiple R-squared:  0.7403, Adjusted R-squared:  0.7393
## F-statistic: 732.8 on 31 and 7968 DF, p-value: < 2.2e-16


linear_model_pred_train = predict.lm(linear_model, newdata = train)
linear_model_pred_test = predict.lm(linear_model, newdata = test)
RMSPE(expm1(linear_model_pred_test), y_test)

## [1] 0.423323

plot(y_test, expm1(linear_model_pred_test),
     main="Stepwise Linear Regression")
abline(a = 0, b=1, col="red", lwd=3, lty=2)

```

Stepwise Linear Regression



*#everything is roughly fine. There is a nice positive correlation
#(even though some predicted price are negative!)*

```
set.seed(1)
dtrain=xgb.DMatrix(as.matrix(x_train),label= log1p(y_train))

best_param = list()
best_seednumber = 1234
best_rmspe = Inf
best_rmspe_index = 0

for (iter in 1:100) {
  param <- list(objective = "reg:linear",
    max_depth = sample(6:20, 1),
    eta = runif(1, .01, .3),
    gamma = runif(1, 0, 2),
    subsample = runif(1, .6, .9),
    colsample_bytree = runif(1, .5, .8),
    min_child_weight = sample(1:40, 1)
  )
  cv.nround = 150
  cv.nfold = 5
  seed.number = sample.int(10000, 1)[[1]]
  set.seed(seed.number)
  mdcv <- xgb.cv(data=dtrain, params = param, nthread=6,
```



```

        nfold=cv.nfold, nrounds=cv.nround,
        verbose = F, early_stopping_rounds=8, maximize=FALSE,
        feval = xg_eval_mae)

min_rmspe = min(mdcv$evaluation_log[, test_error_mean])
min_rmspe_index = which.min(mdcv$evaluation_log[, test_error_mean])

if (min_rmspe < best_rmspe) {
  best_rmspe = min_rmspe
  best_rmspe_index = min_rmspe_index
  best_seednumber = seed.number
  best_param = param
}
}

```

```
## Warning in sqrt(rowMeans(msg^2) - bst_evaluation^2): NaNs produced
```

```
## Warning in sqrt(rowMeans(msg^2) - bst_evaluation^2): NaNs produced
```

```
nround = best_rmspe_index
```

```
set.seed(best_seednumber)
```

```
md <- xgb.train(data=dtrain, params=best_param, nrounds=150, nthread=6)
```

```
# md <- xgb.train(data=dtrain, nrounds=10, nthread=6)
```

```
importance <- xgb.importance(feature_names = names(x_data), model = md)
```

```
importance
```

```
##           Feature      Gain      Cover  Frequency
##  1:              y 0.4479872547 0.2132322296 0.1844444444
##  2:         quality 0.1772495363 0.0947182614 0.0662222222
##  3:              x 0.0983593846 0.1350865891 0.1637777778
##  4:    living_area 0.0813651198 0.0810304353 0.0664444444
##  5:    quality_2 0.0433882807 0.0245767035 0.0188888889
##  6:    land_size 0.0220204466 0.0687431029 0.0988888889
##  7:   bathrooms 0.0179614185 0.0324691180 0.0124444444
##  8:   environment 0.0160290980 0.0454187573 0.0380000000
##  9:    year_built 0.0145541843 0.0366422854 0.0560000000
## 10:    bedrooms 0.0141377252 0.0427980698 0.0244444444
## 11: year_renovation 0.0125296411 0.0459485575 0.0588888889
## 12:    basement 0.0097307429 0.0284452188 0.0244444444
## 13:  maintenance 0.0065748014 0.0272741769 0.0442222222
## 14: environment_2 0.0050339018 0.0145879654 0.0126666667
## 15:  location.f10 0.0047879992 0.0099799043 0.0057777778
## 16: area_per_floor 0.0040088609 0.0133265472 0.0351111111
## 17:  maintenance_2 0.0030541353 0.0152847013 0.0206666667
## 18:  location.f3 0.0028003117 0.0037591124 0.0026666667
## 19:  location.f5 0.0027662759 0.0079873210 0.0066666667
## 20:    floors 0.0025339207 0.0127374021 0.0126666667
## 21:  location.f6 0.0021770329 0.0098396964 0.0062222222
## 22:  location.f12 0.0020967002 0.0075651117 0.0046666667
## 23:  location.f15 0.0018509612 0.0068391564 0.0077777778
## 24:    type.f3 0.0012223868 0.0006491694 0.0004444444
## 25:   has_basement 0.0009864849 0.0023463874 0.0031111111
```

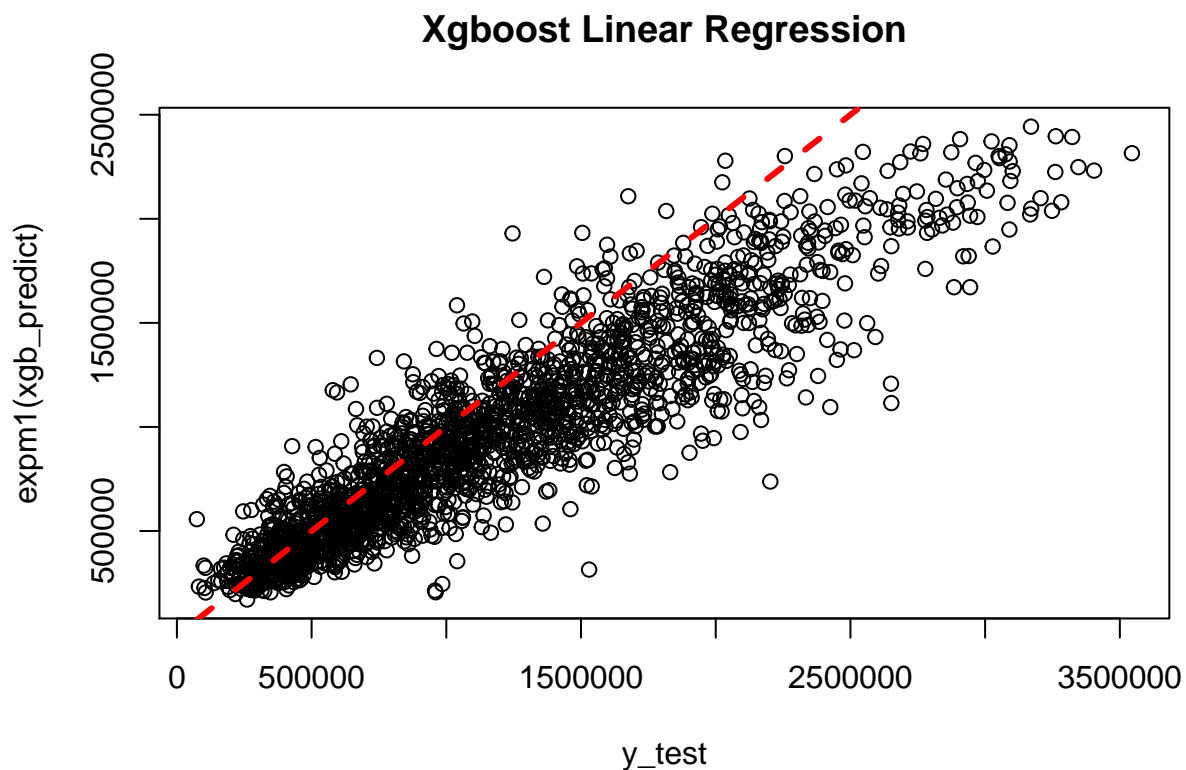
```
## 26:    location.f14 0.0009467910 0.0011941456 0.0015555556
## 27:    location.f8 0.0007438852 0.0034071203 0.0042222222
## 28:    location.f4 0.0006946140 0.0011930131 0.0035555556
## 29:    location.f2 0.0005220565 0.0019119467 0.0028888889
## 30:    location.f7 0.0004291631 0.0017619990 0.0024444444
## 31:    is_renovated 0.0003806944 0.0038771226 0.0024444444
## 32:    location.f11 0.0003188989 0.0009157683 0.0022222222
## 33:    location.f9 0.0002573668 0.0020988151 0.0015555556
## 34:    location.f13 0.0002084679 0.0009603902 0.0015555556
## 35:    type.f2 0.0001748538 0.0011452201 0.0015555556
## 36:    is_old 0.0001166028 0.0002484783 0.0004444444
##      Feature      Gain      Cover      Frequency
```

```
important_features <- head(importance,25)$Feature
```

```
dtest=xgb.DMatrix(as.matrix(x_test))
xgb_predict = predict(md,dtest)
#RMSPE(round_any(xgb_predict,100),y_test)
RMSPE(expm1(xgb_predict),y_test)
```

```
## [1] 0.3223504
```

```
plot(y_test, expm1(xgb_predict),
     main="Xgboost Linear Regression")
abline(a = 0, b=1, col="red", lwd=3, lty=2)
```



Train with the whole dataset

```

ddata=xgb.DMatrix(as.matrix(x_data),label= log1p(y_data))

best_param <- list(
  objective = "reg:linear",
  max_depth = 20,
  eta = 0.03860384,
  gamma = 0.170282,
  subsample = 0.6801184,
  colsample_bytree = 0.7937623,
  min_child_weight = 25
)

best_seednumber <- 6971
best_rmspe_index <- 44

param <- best_param
cv.nround = 150
cv.nfold = 10
set.seed(best_seednumber)
all_mdcv <- xgb.cv(data=ddata, params = param, nthread=6,
  nfold=cv.nfold, nrounds=cv.nround,
  verbose = T, early_stopping_rounds=8, maximize=FALSE,
  feval = xg_eval_mae)

## [1] train-error:0.999998+0.000000 test-error:0.999998+0.000000
## Multiple eval metrics are present. Will use test_error for early stopping.
## Will train until test_error hasn't improved in 8 rounds.
##
## [2] train-error:0.999995+0.000000 test-error:0.999995+0.000000
## [3] train-error:0.999992+0.000000 test-error:0.999992+0.000000
## [4] train-error:0.999987+0.000000 test-error:0.999987+0.000000
## [5] train-error:0.999979+0.000000 test-error:0.999979+0.000000
## [6] train-error:0.999967+0.000000 test-error:0.999967+0.000001
## [7] train-error:0.999951+0.000000 test-error:0.999951+0.000001
## [8] train-error:0.999927+0.000000 test-error:0.999927+0.000002
## [9] train-error:0.999894+0.000000 test-error:0.999894+0.000002
## [10] train-error:0.999848+0.000000 test-error:0.999848+0.000003
## [11] train-error:0.999786+0.000001 test-error:0.999786+0.000004
## [12] train-error:0.999702+0.000001 test-error:0.999702+0.000006
## [13] train-error:0.999591+0.000001 test-error:0.999591+0.000008
## [14] train-error:0.999446+0.000001 test-error:0.999446+0.000011
## [15] train-error:0.999259+0.000002 test-error:0.999258+0.000014
## [16] train-error:0.999019+0.000002 test-error:0.999018+0.000019
## [17] train-error:0.998716+0.000002 test-error:0.998715+0.000023
## [18] train-error:0.998338+0.000003 test-error:0.998336+0.000030
## [19] train-error:0.997868+0.000004 test-error:0.997866+0.000038
## [20] train-error:0.997293+0.000005 test-error:0.997291+0.000047
## [21] train-error:0.996596+0.000005 test-error:0.996592+0.000058
## [22] train-error:0.995755+0.000006 test-error:0.995751+0.000070
## [23] train-error:0.994752+0.000009 test-error:0.994748+0.000084
## [24] train-error:0.993566+0.000010 test-error:0.993560+0.000101
## [25] train-error:0.992173+0.000011 test-error:0.992165+0.000120
## [26] train-error:0.990551+0.000013 test-error:0.990541+0.000143
## [27] train-error:0.988677+0.000016 test-error:0.988664+0.000168

```

## [28]	train-error:0.986523+0.000017	test-error:0.986510+0.000197
## [29]	train-error:0.984071+0.000022	test-error:0.984054+0.000226
## [30]	train-error:0.981294+0.000025	test-error:0.981274+0.000260
## [31]	train-error:0.978167+0.000031	test-error:0.978143+0.000296
## [32]	train-error:0.974672+0.000037	test-error:0.974642+0.000334
## [33]	train-error:0.970787+0.000041	test-error:0.970755+0.000376
## [34]	train-error:0.966493+0.000045	test-error:0.966451+0.000419
## [35]	train-error:0.961771+0.000049	test-error:0.961719+0.000462
## [36]	train-error:0.956608+0.000057	test-error:0.956546+0.000508
## [37]	train-error:0.950994+0.000062	test-error:0.950922+0.000554
## [38]	train-error:0.944917+0.000063	test-error:0.944839+0.000609
## [39]	train-error:0.938371+0.000070	test-error:0.938281+0.000660
## [40]	train-error:0.931347+0.000085	test-error:0.931243+0.000708
## [41]	train-error:0.923859+0.000092	test-error:0.923743+0.000761
## [42]	train-error:0.915892+0.000095	test-error:0.915764+0.000820
## [43]	train-error:0.907457+0.000098	test-error:0.907319+0.000880
## [44]	train-error:0.898565+0.000103	test-error:0.898418+0.000937
## [45]	train-error:0.889223+0.000099	test-error:0.889059+0.000996
## [46]	train-error:0.879433+0.000100	test-error:0.879254+0.001057
## [47]	train-error:0.869244+0.000094	test-error:0.869059+0.001124
## [48]	train-error:0.858638+0.000120	test-error:0.858431+0.001168
## [49]	train-error:0.847634+0.000121	test-error:0.847410+0.001229
## [50]	train-error:0.836289+0.000131	test-error:0.836060+0.001271
## [51]	train-error:0.824583+0.000144	test-error:0.824357+0.001331
## [52]	train-error:0.812564+0.000145	test-error:0.812336+0.001367
## [53]	train-error:0.800263+0.000149	test-error:0.800021+0.001433
## [54]	train-error:0.787695+0.000153	test-error:0.787449+0.001468
## [55]	train-error:0.774879+0.000153	test-error:0.774649+0.001544
## [56]	train-error:0.761886+0.000153	test-error:0.761652+0.001597
## [57]	train-error:0.748700+0.000156	test-error:0.748486+0.001619
## [58]	train-error:0.735355+0.000145	test-error:0.735164+0.001696
## [59]	train-error:0.721911+0.000148	test-error:0.721753+0.001779
## [60]	train-error:0.708362+0.000159	test-error:0.708231+0.001850
## [61]	train-error:0.694768+0.000178	test-error:0.694715+0.001903
## [62]	train-error:0.681124+0.000181	test-error:0.681149+0.002004
## [63]	train-error:0.667467+0.000188	test-error:0.667581+0.002146
## [64]	train-error:0.653823+0.000201	test-error:0.654031+0.002319
## [65]	train-error:0.640249+0.000215	test-error:0.640570+0.002499
## [66]	train-error:0.626735+0.000251	test-error:0.627175+0.002724
## [67]	train-error:0.613310+0.000284	test-error:0.613925+0.002971
## [68]	train-error:0.600015+0.000313	test-error:0.600773+0.003228
## [69]	train-error:0.586845+0.000338	test-error:0.587796+0.003547
## [70]	train-error:0.573844+0.000379	test-error:0.575010+0.003850
## [71]	train-error:0.561020+0.000401	test-error:0.562441+0.004159
## [72]	train-error:0.548390+0.000428	test-error:0.550122+0.004547
## [73]	train-error:0.535981+0.000460	test-error:0.538061+0.004983
## [74]	train-error:0.523799+0.000496	test-error:0.526236+0.005354
## [75]	train-error:0.511839+0.000551	test-error:0.514662+0.005769
## [76]	train-error:0.500145+0.000616	test-error:0.503417+0.006215
## [77]	train-error:0.488741+0.000683	test-error:0.492472+0.006725
## [78]	train-error:0.477598+0.000709	test-error:0.481756+0.007254
## [79]	train-error:0.466735+0.000754	test-error:0.471426+0.007763
## [80]	train-error:0.456182+0.000798	test-error:0.461479+0.008391
## [81]	train-error:0.445890+0.000851	test-error:0.451806+0.008926

```

## [82] train-error:0.435934+0.000912    test-error:0.442456+0.009496
## [83] train-error:0.426271+0.000977    test-error:0.433555+0.010153
## [84] train-error:0.416926+0.000999    test-error:0.424972+0.010822
## [85] train-error:0.407838+0.001069    test-error:0.416677+0.011392
## [86] train-error:0.399112+0.001140    test-error:0.408820+0.012120
## [87] train-error:0.390680+0.001165    test-error:0.401221+0.012679
## [88] train-error:0.382551+0.001220    test-error:0.394065+0.013419
## [89] train-error:0.374719+0.001274    test-error:0.387231+0.014208
## [90] train-error:0.367220+0.001314    test-error:0.380855+0.015028
## [91] train-error:0.360067+0.001334    test-error:0.374733+0.015669
## [92] train-error:0.353133+0.001381    test-error:0.368864+0.016223
## [93] train-error:0.346586+0.001426    test-error:0.363373+0.016880
## [94] train-error:0.340314+0.001486    test-error:0.358276+0.017642
## [95] train-error:0.334280+0.001482    test-error:0.353376+0.018296
## [96] train-error:0.328500+0.001550    test-error:0.348877+0.018849
## [97] train-error:0.323096+0.001583    test-error:0.344738+0.019632
## [98] train-error:0.317859+0.001648    test-error:0.340796+0.020109
## [99] train-error:0.312942+0.001687    test-error:0.337243+0.020907
## [100] train-error:0.308262+0.001780    test-error:0.333889+0.021568
## [101] train-error:0.303735+0.001817    test-error:0.330618+0.022102
## [102] train-error:0.299505+0.001838    test-error:0.327815+0.022851
## [103] train-error:0.295493+0.001932    test-error:0.325082+0.023341
## [104] train-error:0.291698+0.001992    test-error:0.322625+0.023947
## [105] train-error:0.288124+0.002018    test-error:0.320387+0.024404
## [106] train-error:0.284804+0.002083    test-error:0.318439+0.024963
## [107] train-error:0.281635+0.002127    test-error:0.316712+0.025503
## [108] train-error:0.278647+0.002201    test-error:0.315093+0.025959
## [109] train-error:0.275821+0.002312    test-error:0.313724+0.026343
## [110] train-error:0.273148+0.002379    test-error:0.312489+0.026816
## [111] train-error:0.270616+0.002397    test-error:0.311482+0.027413
## [112] train-error:0.268206+0.002384    test-error:0.310615+0.028016
## [113] train-error:0.265961+0.002452    test-error:0.309899+0.028472
## [114] train-error:0.263851+0.002528    test-error:0.309167+0.028785
## [115] train-error:0.261845+0.002589    test-error:0.308606+0.029213
## [116] train-error:0.259964+0.002625    test-error:0.308252+0.029708
## [117] train-error:0.258101+0.002603    test-error:0.307856+0.030168
## [118] train-error:0.256486+0.002565    test-error:0.307593+0.030409
## [119] train-error:0.254953+0.002552    test-error:0.307333+0.030752
## [120] train-error:0.253520+0.002593    test-error:0.307176+0.031080
## [121] train-error:0.252068+0.002508    test-error:0.307123+0.031446
## [122] train-error:0.250691+0.002458    test-error:0.307029+0.031817
## [123] train-error:0.249483+0.002436    test-error:0.307019+0.031937
## [124] train-error:0.248272+0.002407    test-error:0.307076+0.032172
## [125] train-error:0.247101+0.002423    test-error:0.307123+0.032406
## [126] train-error:0.246024+0.002389    test-error:0.307345+0.032636
## [127] train-error:0.244967+0.002394    test-error:0.307646+0.032967
## [128] train-error:0.243939+0.002387    test-error:0.307853+0.033171
## [129] train-error:0.243072+0.002371    test-error:0.308086+0.033347
## [130] train-error:0.242170+0.002355    test-error:0.308292+0.033357
## [131] train-error:0.241313+0.002407    test-error:0.308598+0.033498
## Stopping. Best iteration:
## [123] train-error:0.249483+0.002436    test-error:0.307019+0.031937

```

```

all_best_rmspe = min(all_mdcv$evaluation_log[, test_error_mean])
all_best_rmspe_index = which.min(all_mdcv$evaluation_log[, test_error_mean])
nround = all_best_rmspe_index
set.seed(best_seednumber)
all_md <- xgb.train(data=ddata, params=best_param, nrounds=150, nthread=6)
importance <- xgb.importance(feature_names = names(x_data), model = all_md)
importance

```

##		Feature	Gain	Cover	Frequency
## 1:		y	0.4453769553	0.1827461384	0.1530024168
## 2:		quality	0.1689849655	0.0884615552	0.0817066369
## 3:		x	0.1045758388	0.1397985976	0.1430563302
## 4:		living_area	0.0824370275	0.0820664751	0.0729689533
## 5:		quality_2	0.0328558391	0.0113812619	0.0082729132
## 6:		land_size	0.0253054028	0.0974236222	0.1049451571
## 7:		environment	0.0198923476	0.0493807236	0.0510317903
## 8:		bathrooms	0.0185779152	0.0233716523	0.0109685815
## 9:		bedrooms	0.0148959693	0.0387375233	0.0224948875
## 10:		year_built	0.0145923150	0.0446905106	0.0622792341
## 11:		maintenance	0.0118730750	0.0468660797	0.0676705707
## 12:		year_renovation	0.0111992019	0.0433118890	0.0462911322
## 13:		basement	0.0100211645	0.0330178985	0.0356943670
## 14:		area_per_floor	0.0094144268	0.0312483723	0.0652537646
## 15:		location.f10	0.0055795520	0.0065054110	0.0038111173
## 16:		location.f3	0.0040939363	0.0042329817	0.0025097602
## 17:		environment_2	0.0028742012	0.0089088260	0.0123628927
## 18:		floors	0.0024270269	0.0110713094	0.0097601785
## 19:		location.f12	0.0021790845	0.0029486544	0.0022308979
## 20:		location.f14	0.0020033729	0.0023488911	0.0012084030
## 21:		location.f5	0.0018563061	0.0059619677	0.0035322551
## 22:		location.f15	0.0014255436	0.0040610709	0.0037181632
## 23:		maintenance_2	0.0013745047	0.0042936636	0.0073433724
## 24:		location.f6	0.0010538668	0.0068682196	0.0032533928
## 25:		is_renovated	0.0007342050	0.0099035991	0.0016731735
## 26:		location.f2	0.0007299815	0.0029371081	0.0027886224
## 27:		location.f4	0.0007273756	0.0025945695	0.0030674847
## 28:		location.f8	0.0005451127	0.0034655413	0.0026027143
## 29:		has_basement	0.0005098912	0.0011665556	0.0023238520
## 30:		location.f9	0.0003186369	0.0039457368	0.0021379439
## 31:		location.f13	0.0003023557	0.0014131321	0.0018590816
## 32:		location.f11	0.0003007062	0.0014159545	0.0020449898
## 33:		type.f2	0.0002834165	0.0008303033	0.0027886224
## 34:		location.f7	0.0002519195	0.0010658467	0.0009295408
## 35:		is_old	0.0002168308	0.0006632676	0.0007436326
## 36:		type.f3	0.0002097291	0.0008950906	0.0016731735
##		Feature	Gain	Cover	Frequency

```
important_features <- importance$Feature
```

```

xgb_predict = predict(all_md,dtest)
#RMSPE(round_any(xgb_predict,100),y_test)
all_best_rmspe

```

```
## [1] 0.3070193
```


Lasso Regression

```
library(glmnet)
```

```
## Loading required package: foreach
```

```
## Loaded glmnet 2.0-5
```

```
# lasso_fit_trial <- glmnet(x = as.matrix(x_train), y = y_train, alpha = 1)
# plot(lasso_fit_trial, xvar = "lambda")
crossval <- cv.glmnet(x = as.matrix(x_train), y = log1p(y_train))
plot(crossval)
```



```
penalty <- crossval$lambda.min #optimal lambda
penalty #minimal shrinkage
```

```
## [1] 0.0001064018
```

```
lasso_fit <- glmnet(x = as.matrix(x_train), y = log1p(y_train), alpha = 1, lambda = penalty) #estimate
coef(lasso_fit)
```

```
## 37 x 1 sparse Matrix of class "dgCMatrix"
##              s0
## (Intercept)  1.239180e+01
## x            1.392552e-01
## y            3.980781e-01
## bedrooms     6.540991e-02
## bathrooms    6.192725e-02
## land_size     3.120557e-07
```

```
## living_area      1.278545e-05
## basement        2.506082e-06
## floors           .
## quality          1.915023e+00
## environment      5.133925e-02
## maintenance     -3.493324e-02
## year_built       -1.881513e-03
## year_renovation  1.449444e-03
## is_renovated      2.011766e-01
## has_basement      4.061432e-02
## area_per_floor   -1.907771e-06
## is_old           1.155028e-01
## quality_2        -6.639640e-01
## environment_2     4.987565e-01
## maintenance_2     2.274912e-01
## location.f2       -4.068558e-01
## location.f3       -5.646341e-01
## location.f4       -2.263765e-01
## location.f5       -5.057094e-01
## location.f6        1.898373e-02
## location.f7       -2.402610e-01
## location.f8       -1.564742e-01
## location.f9       -2.205265e-01
## location.f10      -4.194866e-01
## location.f11       6.586722e-03
## location.f12      -2.773149e-01
## location.f13      -1.945964e-01
## location.f14      -8.931695e-02
## location.f15      -3.999123e-01
## type.f2           9.176298e-02
## type.f3           8.076675e-02
```

```
lasso_predict = predict(lasso_fit, newx = as.matrix(x_test), s = penalty)
RMSPE(expm1(lasso_predict), y_test)
```

```
## [1] 0.4233612
```

After comparison, XGBOOST linear regression has the lowest RMSPE. Final output is made based on the regression model with best parameters.

```
#let's create a submission file
```

```
doutput=xgb.DMatrix(as.matrix(data_output))
output = expm1(predict(md,doutput))
submission = data.frame(Id = output_id,
                        Prediction = output)

#submission_linear_model = data.frame(Id = data_output[, "Id"],
#                                     Prediction = linear_model_pred_test)
#let us create the submission file
write.table(x = submission,
            file = "prediction_8_md.csv",
            sep = ",", row.names = FALSE, col.names = TRUE)
```